



Video: In a factory setting, a puppet is talking to another puppet. The camera captures a close-up of the speaking puppet's face.



Audio: The puppet's speaking voice, ambient sounds of the factory, and atmospheric background music.

Human eval: [1.0, 1.0, 0.8] -> 0.93
AVC-RewardModel (ours): [1.0, 1.0, 0.9] -> 0.97

Qwen3-Omni-Thinking (baseline): 0.46
Qwen3-Omni-Instruct: 0.69
Gemini3.1pro: 0.76 | AV-Align: 0.28
Desync↓: 0.0 | ImageBind: 0.32 | PEAVS: 3.44
JavisScore: 0.25 | AVHScore: 0.27

"Overall, the audio-video content is highly consistent. The macro-level atmosphere and thematic elements are perfectly aligned. The fine-grained synchronization of the speech is excellent. The only minor discrepancy is the lack of a corresponding visual for the ambient **background music and drone**. This is a **common convention in film and animation and does not significantly detract from the overall coherence** or quality of the production."



Video: Two girls are having a conversation in a classroom. The girl on the left speaks first, and then the girl on the right replies to her.



Audio: The first half features clear female speaking voices, while the second half contains noise resembling birdsong, making the speech unintelligible.

Human eval: [0.7, 0.3, 0.5] -> 0.50
AVC-RewardModel (ours): [0.4, 0.4, 0.4] -> 0.40

Qwen3-Omni-Thinking (baseline): 0.10
Qwen3-Omni-Instruct: 1.00
Gemini3.1pro: 0.10 | AV-Align: 0.00
Desync↓: 2.0 | ImageBind: 0.23 | PEAVS: 4.93
JavisScore: 0.20 | AVHScore: 0.25

"The audio and video content are highly inconsistent and appear to be from two completely different and unrelated scenes. There is a severe mismatch in the overall theme, environment, character actions, and emotional context. **The only point of minor connection is the presence of a feminine-like voice that matches one of the two anime characters on screen**, but this single element is not enough to overcome the fundamental contradictions between the two streams."



Video: A musical performance in a concert hall. The first few frames show the performance, and the latter half shows the orchestra packing up and taking a curtain call at the end.



Audio: The music played by the orchestra (violins, etc.) is heard throughout the entire duration.

Human eval: [0.7, 0.7, 0.7] -> 0.70
AVC-RewardModel (ours): [1.0, 0.7, 0.7] -> 0.80

Qwen3-Omni-Thinking (baseline): 0.92
Qwen3-Omni-Instruct: 1.00
Gemini3.1pro: 0.10 | AV-Align: 0.25
Desync↓: 1.6 | ImageBind: 0.44 | PEAVS: 5.0
JavisScore: 0.31 | AVHScore: 0.35

"Overall, the video presents a significant inconsistency. While it is thematically coherent at a high level, it fails critically in terms of temporal synchronicity and fine-grained event correspondence. **The musicians are miming to a studio recording, which breaks the authentic physical connection between their actions and the sound produced**. This results in a performance that feels artificial and disconnected."



Video: With a fireplace in the background, a worker is hammering high-temperature iron, striking it twice in total.



Audio: Faint background sound of a burning fire, and the sound of a hammer striking the iron. There are two strikes in total, but they are delayed compared to the striking moments shown in the video.

Human eval: [1.0, 0.2, 0.3] -> 0.5
AVC-RewardModel (ours): [1.0, 0.4, 0.4] -> 0.6

Qwen3-Omni-Thinking (baseline): 0.92
Qwen3-Omni-Instruct: 0.92
Gemini3.1pro: 0.67 | AV-Align: 0.11
Desync↓: 2.0 | ImageBind: 0.28 | PEAVS: 3.67
JavisScore: 0.18 | AVHScore: 0.20

"Overall, the audio-video content exhibits a high degree of thematic consistency but suffers from a critical failure in temporal and event-level synchronization. While the general atmosphere is cohesive, **the lack of precise timing between the hammer** impacts and the resulting sounds makes the experience feel artificial and poorly produced. The audio is clearly dubbed and not synchronized with the visual actions."